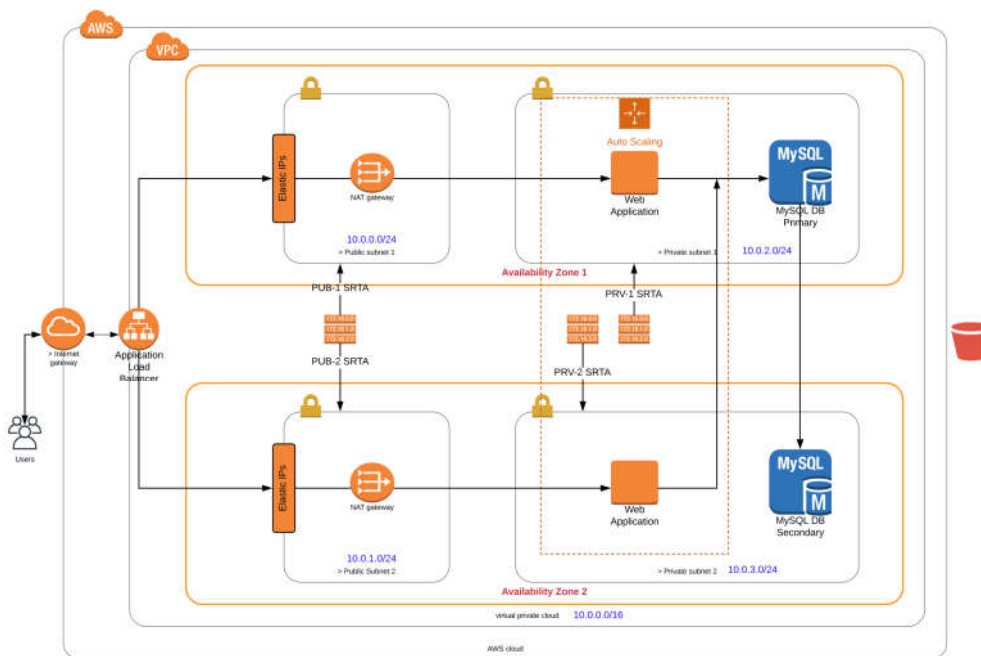


## 3.C. Networking Infrastructure

- Network Implementation Diagram
  - Adding Subnets
  - Adding a NAT Gateway
  - Routing
    - Route Tables
    - Routes
    - SubnetRouteTableAssociation
  - Outputs
    - Join Function

### Network Implementation Diagram



#### Don't hard-code parameters

Avoid hard coding parameter values. Instead, use a separate parameter file to store parameter values. Note that the parameter file should be in .json format, as .yaml format is not yet supported for the parameter file.

#### Location Dependencies

- `Parameters` should be declared above your `Resources`.
- CloudFormation knows to create the resources in order, based on their dependencies (VPC and InternetGateway, before creating the InternetGatewayAttachment).

#### Default Parameters

You can also provide default values for parameters in case one was not passed in. In this example you can see that `VpcCIDR` has a default value of `10.0.0.0/16`.

## Adding Subnets

To specify a `Subnet` for your `VPC` you use the following syntax:

```
1 Type: AWS::EC2::Subnet
2 Properties:
3   AssignIpv6AddressOnCreation: Boolean
4   AvailabilityZone: String
5   CidrBlock: String
6   Ipv6CidrBlock: String
7   MapPublicIpOnLaunch: Boolean
8   Tags:
9     - Tag
10  VpcId: String
```

Here is the actual setup of our 2 private `Subnets` :

```
1 PrivateSubnet1:
2   Type: AWS::EC2::Subnet
3   Properties:
4     VpcId: !Ref VPC
5     AvailabilityZone: !Select [ 0, !GetAZs '' ]
6     CidrBlock: !Ref PrivateSubnet1CIDR
7     MapPublicIpOnLaunch: false
8     Tags:
9       - Key: Name
10       Value: !Sub ${EnvironmentName} Private Subnet (AZ1)
11
12 PrivateSubnet2:
13   Type: AWS::EC2::Subnet
14   Properties:
15     VpcId: !Ref VPC
16     AvailabilityZone: !Select [ 1, !GetAZs '' ]
17     CidrBlock: !Ref PrivateSubnet2CIDR
18     MapPublicIpOnLaunch: false
19     Tags:
20       - Key: Name
21       Value: !Sub ${EnvironmentName} Private Subnet (AZ2)
```

You can see the index being used from the returning `AvailabilityZone`'s array. Notice that our `subnets` are not sharing `AvailabilityZones`. We are keeping them separated like we displayed in our diagrams from the previous lesson:

PrivateSubnet1: `AvailabilityZone: !Select [ 0, !GetAZ's '' ]`

PrivateSubnet2: `AvailabilityZone: !Select [ 1, !GetAZ's '' ]`

This code:

```
1 !select [0, !GetAZs'']
```

calls the function `GetAZ`, which returns a list of availability zones, which are indexed 0, 1 etc.

#### Tip

Name your subnets using tags, to keep track when you create many subnets.

## Adding a NAT Gateway

You can use `NAT Gateways` in both your public and/or private `Subnets`. The following code is the basic syntax for declaring a `NAT Gateway`:

```
1 Type: AWS::EC2::NatGateway
2 Properties:
3   AllocationId: String
4   SubnetId: String
5   Tags:
6     - Tag
```

The following declarations are from the sample code shown in the above video:

```
1 NatGateway1EIP:
2   Type: AWS::EC2::EIP
3   DependsOn: InternetGatewayAttachment
4   Properties:
5     Domain: vpc
6
7 NatGateway2EIP:
8   Type: AWS::EC2::EIP
9   DependsOn: InternetGatewayAttachment
10  Properties:
11    Domain: vpc
12
13 NatGateway1:
14   Type: AWS::EC2::NatGateway
15   Properties:
16     AllocationId: !GetAtt NatGateway1EIP.AllocationId
17     SubnetId: !Ref PublicSubnet1
18
19 NatGateway2:
20   Type: AWS::EC2::NatGateway
21   Properties:
22     AllocationId: !GetAtt NatGateway2EIP.AllocationId
23     SubnetId: !Ref PublicSubnet2
```

The `EIP` in `AWS::EC2::EIP` stands for Elastic IP. This will give us a known/constant IP address to use instead of a disposable or ever-changing IP address. This is important when you have applications that depend on a particular IP address. `NatGateway1EIP` uses this type for that very reason:

```
1 NatGateway1EIP:
```

```

2 |     Type: AWS::EC2::EIP
3 |     DependsOn: InternetGatewayAttachment
4 |     Properties:
5 |         Domain: vpc

```

Tip:

- Use the `DependsOn` attribute to protect your dependencies from being created without the proper requirements.

In the scenario above the `EIP` allocation will only happen after the `InternetGatewayAttachment` has completed.

## Routing

**Routing:** Routing is the action of applying routing rules to your network, in this case, to your VPC.

**Routing rule:** Resources follow the routing rule, which defines what resource has access to communicate with another resource. It blocks traffic from resources that do not follow the routing rule.

### Route Tables

We create `RouteTables` for `VPCs` so that we can add `Routes` that we later associate with `Subnets`. The following is the syntax used to define a `RouteTable`:

```

1 | Type: AWS::EC2::RouteTable
2 | Properties:
3 |     Tags:
4 |         - Tag
5 |     VpcId: String

```

The only required property for setting up a `RouteTable` is the `VpcId`. Here is an example table from the video lesson:

```

1 | PublicRouteTable:
2 |     Type: AWS::EC2::RouteTable
3 |     Properties:
4 |         VpcId: !Ref VPC
5 |         Tags:
6 |             - Key: Name
7 |               Value: !Sub ${EnvironmentName} Public Routes

```

### Routes

The following is the syntax used to set up our `Route`:

```

1 | Type: AWS::EC2::Route
2 | Properties:
3 |     DestinationCidrBlock: String
4 |     DestinationIpv6CidrBlock: String
5 |     EgressOnlyInternetGatewayId: String

```

```

6 | GatewayId: String
7 | InstanceId: String
8 | NatGatewayId: String
9 | NetworkInterfaceId: String
10 | RouteTableId: String
11 | VpcPeeringConnectionId: String

```

The `DestinationCidrBlock` property is used for destination matching and a wildcard address (`0.0.0.0/0`) to reference all traffic.

So in the following example, when we use the wildcard address `0.0.0.0/0`, we are saying for any address that comes through this route, send it to the referenced `GatewayId`:

```

1 | DefaultPublicRoute:
2 |     Type: AWS::EC2::Route
3 |     DependsOn: InternetGatewayAttachment
4 |     Properties:
5 |         RouteTableId: !Ref PublicRouteTable
6 |         DestinationCidrBlock: 0.0.0.0/0
7 |         GatewayId: !Ref InternetGateway

```

### SubnetRouteTableAssociation

In order to associate `Subnets` with our `Route Table` we will need to use a `SubnetRouteTableAssociation`. `SubnetRouteTableAssociations` are defined using the following syntax:

```

1 | Type: AWS::EC2::SubnetRouteTableAssociation
2 | Properties:
3 |     RouteTableId: String
4 |     SubnetId: String

```

This only takes two properties, which are the id's used for our `RouteTable` and our `Subnet`. You can see references used in the example from our video lesson above:

```

1 | PublicSubnet1RouteTableAssociation:
2 |     Type: AWS::EC2::SubnetRouteTableAssociation
3 |     Properties:
4 |         RouteTableId: !Ref PublicRouteTable
5 |         SubnetId: !Ref PublicSubnet1

```

Important Note:

- `Routes` should be defined starting with the most specific rule and transitioning to the least specific rule.
- Private-VPC servers are not going to have public IP address.
  - So, even if they are placed on a public subnet there is no way to access them if they have no IP address.
- Having 2 Route-Tables to cover both private-subnets seem unnecessary, however this is a good practice in order to take into account future network expansion.

```
1 PublicRouteTable:
2   Type: AWS::EC2::RouteTable
3   Properties:
4     VpcId: !Ref VPC
5     Tags:
6       - Key: Name
7         Value: !Sub ${EnvironmentName} Public Routes
8
9 DefaultPublicRoute:
10  Type: AWS::EC2::Route
11  DependsOn: InternetGatewayAttachment
12  Properties:
13    RouteTableId: !Ref PublicRouteTable
14    DestinationCidrBlock: 0.0.0.0/0
15    GatewayId: !Ref InternetGateway
16
17 PublicSubnet1RouteTableAssociation:
18  Type: AWS::EC2::SubnetRouteTableAssociation
19  Properties:
20    RouteTableId: !Ref PublicRouteTable
21    SubnetId: !Ref PublicSubnet1
22
23 PublicSubnet2RouteTableAssociation:
24  Type: AWS::EC2::SubnetRouteTableAssociation
25  Properties:
26    RouteTableId: !Ref PublicRouteTable
27    SubnetId: !Ref PublicSubnet2
28
29
30 PrivateRouteTable1:
31  Type: AWS::EC2::RouteTable
32  Properties:
33    VpcId: !Ref VPC
34    Tags:
35      - Key: Name
36        Value: !Sub ${EnvironmentName} Private Routes (AZ1)
37
38 DefaultPrivateRoute1:
39  Type: AWS::EC2::Route
40  Properties:
41    RouteTableId: !Ref PrivateRouteTable1
42    DestinationCidrBlock: 0.0.0.0/0
43    NatGatewayId: !Ref NatGateway1
44
45 PrivateSubnet1RouteTableAssociation:
46  Type: AWS::EC2::SubnetRouteTableAssociation
47  Properties:
48    RouteTableId: !Ref PrivateRouteTable1
49    SubnetId: !Ref PrivateSubnet1
50
51 PrivateRouteTable2:
52  Type: AWS::EC2::RouteTable
53  Properties:
54    VpcId: !Ref VPC
```

```

55     Tags:
56         - Key: Name
57           Value: !Sub ${EnvironmentName} Private Routes (AZ2)
58
59 DefaultPrivateRoute2:
60     Type: AWS::EC2::Route
61     Properties:
62         RouteTableId: !Ref PrivateRouteTable2
63         DestinationCidrBlock: 0.0.0.0/0
64         NatGatewayId: !Ref NatGateway2
65
66 PrivateSubnet2RouteTableAssociation:
67     Type: AWS::EC2::SubnetRouteTableAssociation
68     Properties:
69         RouteTableId: !Ref PrivateRouteTable2
70         SubnetId: !Ref PrivateSubnet2

```

If your servers have no internet access it's probably because...

- You created the internet gateway but forgot to attach it to your VPC
- You placed your NAT Gateways inside private subnets with no routes to the outside world
- You have a missing route in your routing table
- You created a routing table but forgot to associate your subnet(s) with it.

Documentation:

- [Route Tables Overview](#)
- [Route Table Documentation](#)
- [Route Documentation](#)
- [Subnet Route Table Association Documentation](#)

## Outputs

`Outputs` are optional but are very useful if there are output values you need to:

- import into another stack
- return in a response
- view in AWS console

To declare an `Output` use the following syntax:

```

1  Outputs:
2      Logical ID:
3          Description: Information about the value
4          Value: Value to return
5          Export:
6              Name: Value to export

```

The `Value` is required but the `Name` is optional. In the following example we are returning the id of our `VPC` as well as our Environment's Name:

```
1 VPC:
2     Description: A reference to the created VPC
3     Value: !Ref VPC
4     Export:
5     Name: !Sub ${EnvironmentName}-VPCID
```

## Join Function

You can use the `join` function to combine a group of `values`. The syntax requires you provide a `delimiter` and a list of values you want appended.

`Join` function syntax:

```
1 Fn::Join: [ delimiter, [ comma-delimited list of values ] ]
```

In the following example we are using `!Join` to combine our subnets before returning their values:

```
1 PublicSubnets:
2     Description: A list of the public subnets
3     Value: !Join [ ",", [ !Ref PublicSubnet1, !Ref PublicSubnet2 ] ]
4     Export:
5     Name: !Sub ${EnvironmentName}-PUB-NETS
```

Documentation:

- [Outputs Documentation](#)
- [Join Function](#)
- [Substitutes](#)

[Top](#)

## Deleting VPC

Are you sure that you want to delete this VPC (vpc-084aa83cfa00ca370 UdacityProject)?

Deleting this VPC will also delete these objects associated with this VPC in this region:

- Subnets
- Security Groups
- Network ACLs
- Internet Gateways
- Egress Only Internet Gateways
- Route Tables
- Network Interfaces
- Peering Connections
- Endpoints